

```
Activities Terminal Aug 11 11:35
hadoop@abi-cse128: ~/WordCountTutorial

hadoop@abi-cse128:~$ hadoop version
Hadoop 3.2.3
Source code repository https://github.com/apache/hadoop -r abe358143720085498613d399be3bbf01e0f131
Compiled by ubuntu on 2022-03-20T01:18Z
Compiled with protoc 2.5.0
From source with checksum 39bb14faec14b3aa25388a6d7c345fe8
This command was run using /home/hadoop/hadoop-3.2.3/share/hadoop/common/hadoop-common-3.2.3.jar
hadoop@abi-cse128:~$ javac -version
javac 1.8.0_312
hadoop@abi-cse128:~$ export HADOOP_CLASSPATH=$(hadoop classpath)
hadoop@abi-cse128:~$ echo $HADOOP_CLASSPATH
/home/hadoop/hadoop-3.2.3/etc/hadoop:/home/hadoop/hadoop-3.2.3/share/hadoop/common/lib/*:/home/hadoop/hadoop-3.2.3/share/hadoop/common/*:/home/hadoop/hadoop-3.2.3/share/hadoop/hdfs/lib/*:/home/hadoop/hadoop-3.2.3/share/hadoop/hdfs/*:/home/hadoop/hadoop-3.2.3/share/hadoop/mapreduce/lib/*:/home/hadoop/hadoop-3.2.3/share/hadoop/mapreduce/*:/home/hadoop/hadoop-3.2.3/share/hadoop/yarn/lib/*:/home/hadoop/hadoop-3.2.3/share/hadoop/yarn/*
hadoop@abi-cse128:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [abi-cse128]
Starting resourcemanager
Starting nodenaggers
hadoop@abi-cse128:~$ hadoop fs -mkdir hdfs://localhost:9000/wordCountTutorial
mkdir: 'hdfs://localhost:9000/wordCountTutorial': File exists
hadoop@abi-cse128:~$ hadoop fs -mkdir hdfs://localhost:9000/wordCountTutorial/input
mkdir: 'hdfs://localhost:9000/wordCountTutorial/input': File exists
hadoop@abi-cse128:~$ hadoop fs -put '/home/hadoop/input_data/input.txt' /wordCountTutorial/input
put: '/home/hadoop/input_data/input.txt': No such file or directory
hadoop@abi-cse128:~$ hadoop fs -put '/home/hadoop/Desktop/WordCountTutorial/input.txt' /wordCountTutorial/input
put: '/home/hadoop/Desktop/WordCountTutorial/input.txt': No such file or directory
hadoop@abi-cse128:~$ hadoop fs -put '/home/hadoop/WordCountTutorial/input_data/input.txt' /wordCountTutorial/input
hadoop@abi-cse128:~$ ls
dc Desktop Documents hadoop-3.2.3.tar.gz javac.activation-api-1.2.0.jar.2 Public tmpdata
Downloads javac.activation-api-1.2.0.jar Music Templates Videos
dfsdata hadoop-3.2.3 javac.activation-api-1.2.0.jar.1 Pictures test WordCountTutorial
hadoop@abi-cse128:~$ javac -classpath $(HADOOP_CLASSPATH) -d '/home/hadoop/WordCountTutorial/tutorial_classes' '/home/hadoop/WordCountTutorial/WordCount.java'
```

```
Activities Terminal Aug 11 11:36
hadoop@abi-cse128: ~/WordCountTutorial

dfsdata hadoop-3.2.3 javac.activation-api-1.2.0.jar.1 Pictures test WordCountTutorial
hadoop@abi-cse128:~$ javac -classpath $(HADOOP_CLASSPATH) -d '/home/hadoop/WordCountTutorial/tutorial_classes' '/home/hadoop/WordCountTutorial/WordCount.java'
hadoop@abi-cse128:~/WordCountTutorial$ cd /home/hadoop/WordCountTutorial
hadoop@abi-cse128:~/WordCountTutorial$ javac -classpath $(HADOOP_CLASSPATH) -d '/home/hadoop/WordCountTutorial/tutorial_classes' '/home/hadoop/WordCountTutorial/WordCount.java'
hadoop@abi-cse128:~/WordCountTutorial$ jar -cvf firstTutorial.jar -C tutorial_classes/ .
added manifest
adding: WordCount$TokenizerMapper.class (in = 1736) (out= 754)(deflated 56%)
adding: WordCount.class (in = 1491) (out= 814)(deflated 45%)
adding: WordCount$IntSumReducer.class (in = 1739) (out= 739)(deflated 57%)
hadoop@abi-cse128:~/WordCountTutorial$ hadoop jar '/home/hadoop/WordCountTutorial/firstTutorial.jar' WordCount /wordCountTutorial/input /wordCountTutorial/output
hadoop: command not found or invalid user 'hadoop'
hadoop@abi-cse128:~/WordCountTutorial$ hadoop jar '/home/hadoop/WordCountTutorial/firstTutorial.jar' WordCount /wordCountTutorial/input /wordCountTutorial/output
hadoop: command not found or invalid user 'hadoop'
hadoop@abi-cse128:~/WordCountTutorial$ hadoop jar '/home/hadoop/WordCountTutorial/firstTutorial.jar' WordCount /wordCountTutorial/input /wordCountTutorial/output
2022-08-11 11:31:51,522 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-08-11 11:31:51,775 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-08-11 11:31:51,824 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_160196831543_0001
2022-08-11 11:31:52,055 INFO input.FileInputFormat: Total input files to process : 1
2022-08-11 11:31:52,657 INFO mapreduce.JobSubmitter: number of splits:1
2022-08-11 11:31:52,999 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_160196831543_0001
2022-08-11 11:31:53,000 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-08-11 11:31:53,147 INFO conf.Configuration: resource-types.xml not found
2022-08-11 11:31:53,149 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-08-11 11:31:53,545 INFO impl.YarnClientImpl: Submitted application application_160196831543_0001
2022-08-11 11:31:53,588 INFO mapreduce.Job: The url to track the job: http://abi-cse128:8088/proxy/application_160196831543_0001/
2022-08-11 11:31:53,589 INFO mapreduce.Job: Running job: job_160196831543_0001
2022-08-11 11:31:59,070 INFO mapreduce.Job: Job job_160196831543_0001 running in uber mode : false
2022-08-11 11:32:03,672 INFO mapreduce.Job: map 0% reduce 0%
2022-08-11 11:32:03,736 INFO mapreduce.Job: map 100% reduce 0%
2022-08-11 11:32:07,764 INFO mapreduce.Job: map 100% reduce 100%
2022-08-11 11:32:08,783 INFO mapreduce.Job: Job job_160196831543_0001 completed successfully
```

```
Activities Terminal Aug 11 11:36
hadoop@abi-cse128: ~/WordCountTutorial

Total committed heap usage (bytes)=298844160
Peak Map Physical memory (bytes)=25187208
Peak Map Virtual memory (bytes)=2719883264
Peak Reduce Physical memory (bytes)=176910336
Peak Reduce Virtual memory (bytes)=2732421120

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=94

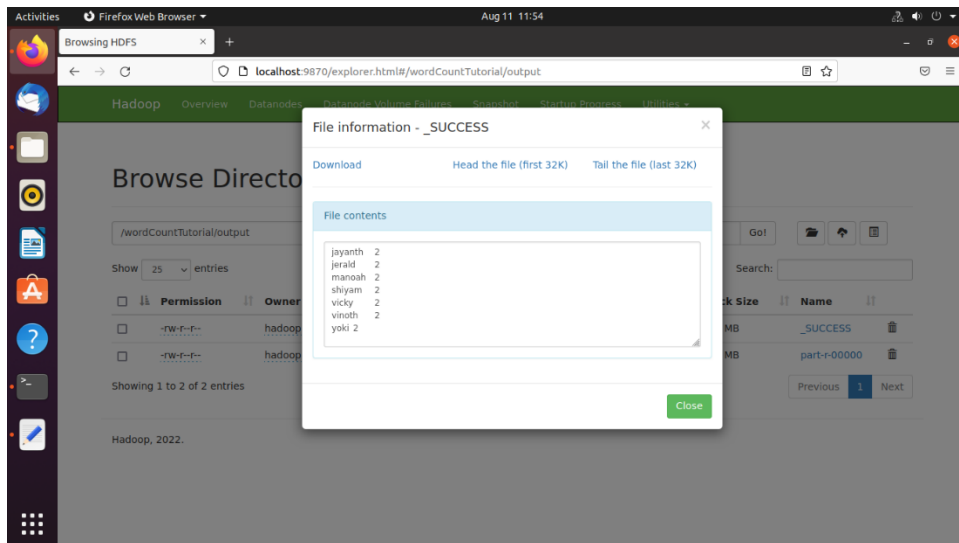
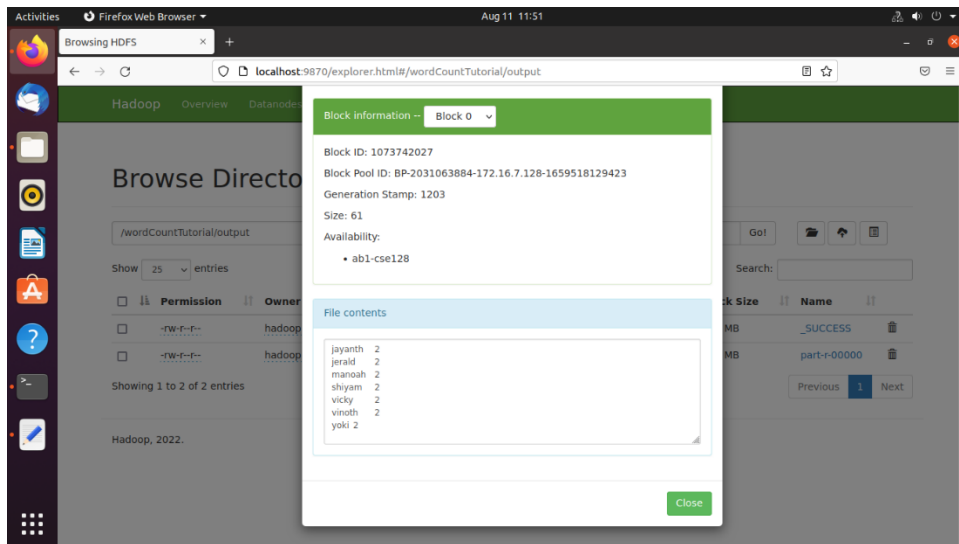
File Output Format Counters
Bytes Written=61
hadoop@abi-cse128:~/WordCountTutorial$ hadoop dfs -cat /wordCountTutorial/output/*
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.
jayanth 2
jerald 2
manoah 2
shiyam 2
vicky 2
vinoth 2
yokl 2
hadoop@abi-cse128:~/WordCountTutorial$ hadoop fs -cat hdfs://localhost:9000/wordCountTutorial/output/*
jayanth 2
jerald 2
manoah 2
shiyam 2
vicky 2
vinoth 2
yokl 2
hadoop@abi-cse128:~/WordCountTutorial$ hadoop fs -cat hdfs://localhost:9000/wordCountTutorial/output
cat: 'hdfs://localhost:9000/wordCountTutorial/output': is a directory
hadoop@abi-cse128:~/WordCountTutorial$
```

```
Activities Terminal Aug 11 11:42 hadoop@ab1-cse128: ~/WordCountTutorial

gridmix
jar <jar> submit a mix of synthetic job, modeling a profiled from production load
jniptath prints the java.library.path
kdial Diagnose Kerberos Problems
kerbname show auth to local principal conversion
key manage keys via the keyProvider
rumenfolder scale a rumen input trace
rumentrace convert logs into a rumen trace
s3guard manage metadata on S3
trace view and modify Hadoop tracing settings
version print the version

Daemon Commands:
kms run KMS, the Key Management Server

SUBCOMMAND may print help when invoked w/o parameters or with -h.
hadoop@ab1-cse128:~/WordCountTutorial$ hadoop dfs -ls /wordCountTutorial/output/*
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement 'hdfs dfs' instead.
-rw-r--r-- 1 hadoop supergroup 0 2022-08-11 11:32 /wordCountTutorial/output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 61 2022-08-11 11:32 /wordCountTutorial/output/part-r-00000
hadoop@ab1-cse128:~/WordCountTutorial$ hadoop dfs -cat /wordCountTutorial/output/_SUCCESS
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement 'hdfs dfs' instead.
hadoop@ab1-cse128:~/WordCountTutorial$ hadoop dfs -cat /wordCountTutorial/output/part-r-00000
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement 'hdfs dfs' instead.
jayanth 2
jerald 2
manoah 2
shiyam 2
vicky 2
vinoth 2
yoki 2
```



```

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {
    public static class TokenizerMapper extends Mapper<Object, Text, Text,
IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context) throws
IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer extends
Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();
        public void reduce(Text key, Iterable<IntWritable> values, Context
context) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);

```

```

        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

COMMANDS

hadoop version

javac -version

export HADOOP_CLASSPATH=\$(hadoop classpath)

echo \$HADOOP_CLASSPATH

hadoop fs -mkdir /wordCountTutorial

start-all.sh

hadoop fs -mkdir hdfs://localhost:9000/wordCountTutorial

hadoop fs -mkdir hdfs://localhost:9000/wordCountTutorial/input

hadoop fs -put '/home/hadoop/input_data/input.txt' /wordCountTutorial/input

ls

javac -classpath \${HADOOP_CLASSPATH} -d '/home/hadoop/tutorial_classes'
 '/home/hadoop/WordCount.java'

jar -cvf firstTutorial.jar -C tutorial_classes/ .

hadoop jar '/home/hadoop/firstTutorial.jar' WordCount /wordCountTutorial/input
 /wordCountTutorial/output

hadoop dfs -cat /wordCountTutorial/output/*

hadoop fs -cat hdfs://localhost:9000/wordCountTutorial/output/*

hadoop fs -cat hdfs://localhost:9000/wordCountTutorial/output

hadoop dfs -cat /wordCountTutorial/output/*

hadoop hdfs -cat /wordCountTutorial/output/*

hadoop dfs -ls /wordCountTutorial/output/*

hadoop dfs -cat /wordCountTutorial/output/_SUCCESS

hadoop dfs -cat /wordCountTutorial/output/part-r-00000